

Global threshold recalibration

One threshold per fixed score, calibrated on all 36 development observations; evaluated on 72 observations from four reused test scenes. All branch times and arm conditions pooled. Scores and fusion weights unchanged.

Model	Old BA	New BA	Successes accepted (new)	Incomplete accepted (new)
DT	0.713	0.849	33/43	2/29
S1	0.654	0.663	14/43	0/29
S3	0.654	0.734	29/43	6/29
S5	0.654	0.711	27/43	6/29
DT + S1	0.654	0.820	29/43	1/29
DT + S3	0.654	0.767	23/43	0/29
DT + S5	0.654	0.767	23/43	0/29
DINov2	0.654	0.861	34/43	2/29
Robometer progress	N/A	0.500	0/43	0/29
Robometer success	0.512	0.706	40/43	15/29
Elapsed time	0.496	0.500	0/43	0/29
Always successful	0.500	0.500	43/43	29/29

Threshold selection: One global threshold per method maximizes pooled DEVELOPMENT balanced accuracy over all six conditions. Ties: fewer false successes, then higher threshold.

Limits:

- Four reused test scenes; post-hoc diagnostic, not a fresh test benchmark.
- Good within-condition AUR0C does not imply pooled score separability.
- 90% cells contain only successes: no cell balanced accuracy or false-success rate.
- Matched visual goals for PSI/DINO; language and longer causal prefixes for Robometer.
- Moving arms to goal pose also adds ten seconds of simulation.

Protocol

```
{
  "input": "/ccn2b/dataset/brd/rwm_progress_tracking/synthetic_mugs_native256_v4/predictions.csv",
  "input_sha256": "c2700e8fcb0c0feaa4d7c3214e4c66a3abe24c2cb0acf6f56ff6ca2de568c64",
  "source_groups": {
    "dev": [
      "episode_019e298a-84a1-7557-ae97-93326347b45c",
      "episode_019e3c90-fb67-765e-8e8e-12c56ddda3ba"
    ],
    "test": [
      "episode_019e1e35-af9f-7d66-93b0-478d1c079799",
      "episode_019e2261-5abe-7ece-910d-8988dd12106a",
      "episode_019e2d91-6e08-7f79-a3e4-c843c73a3ef2",
      "episode_019e3c25-a92e-7824-a264-f1da48b9bd3d"
    ]
  },
  "development_observations_per_method": 36,
  "test_observations_per_method": 72,
  "selection": "One global threshold per method maximizes pooled DEVELOPMENT balanced accuracy over all six conditions. Ties: fewer false successes, then higher threshold.",
  "weighting": "Every development observation equally weighted within its outcome class; same six conditions and three modes per source. All-success 90% conditions participate in calibration and pooled testing.",
  "prediction": "Successful if fixed score >= threshold. Score direction, source split, labels, features and fusion weights unchanged.",
  "primary": "Pooled test balanced accuracy, success recall and false-success rate; report all six cells separately in downloadable results.",
  "controls": "Elapsed time calibrated identically; Always successful remains constant at
```

```
threshold 0.5.",
"robometer": "Both heads calibrated identically. Success old threshold = native 0.5; progress
has no old threshold and is newly calibrated here.",
"uncertainty": "2000 paired test-source cluster bootstrap draws, thresholds fixed. Does not
capture calibration uncertainty from only two dev sources.",
"limitations": [
  "Four reused test scenes; post-hoc diagnostic, not a fresh test benchmark.",
  "Good within-condition AUROC does not imply pooled score separability.",
  "90% cells contain only successes: no cell balanced accuracy or false-success rate.",
  "Matched visual goals for PSI/DINO; language and longer causal prefixes for Robometer.",
  "Moving arms to goal pose also adds ten seconds of simulation."
]
}
```